

How to design a database?

Heli Helskyaho, Matias Helskyaho

KCDC 2024



Titanium Sponsors



Platinum Sponsors

Gold Sponsors



Friends of KCDC







Introduction, Heli



- * Graduated from University of Helsinki (Master of Science, computer science), currently a doctoral student, researcher and lecturer (databases, Multi-model/Converged Databases) at University of Helsinki
- * Worked with Oracle products since 1993, worked for IT since 1990
- * Data and Database!
- * CEO for Miracle Finland Oy
- * Oracle ACE Director
- * Public speaker and an author
- * Author of the book Oracle SQL Developer Data Modeler for Database Design Mastery (Oracle Press, 2015), co-author for Real World SQL and PL/SQL: Advice from the Experts (Oracle Press, 2016), Machine Learning for Oracle Database Professionals: Deploying Model-Driven Applications and Automation Pipelines (Apress, 2021), and Extending Oracle Application Express with Oracle Cloud Features: A Guide to Enhancing APEX Web Applications with Cloud-Native and Machine Learning Technologies (Apress, 2022)



Books



Matias



Oracle ACE
Pro

- * Consultant
 - * Miracle Finland Oy
- * Who am I?
 - * On IT since birth
 - * Professionally a couple of years
 - * OCI, networks, IOT, ML, analytics,...
- * Hobbies
 - * Love learning cool stuff
 - * Playing with tech devices



Why to design?

- * "Data is the most valuable property in our company"
- * "Why do we need to design the database? We already design the application!"

Why is designing the application not enough?

- * Point of view (saving and retrieving data vs. UI)
- * First increment vs. 20 years from now
- * "the whole picture" vs. increments
- * Different goals/targets:
 - * Code tables vs. Code files (how about the data integrity?)
 - * How about analysis, reports, ... everything else but the UI that the data is used for
- * Same terminology, different meaning -> misunderstandings
- * ...

Why to model the data?

- * To facilitate communication about the requirements
- * To find the questions that should be asked
- * To understand the requirements

What is database design?

- * 4 (5) phases, over and over again
 - * Requirement analysis (DM: logical)
 - * Conceptual design (DM: logical)
 - * Logical design (DM: relational)
 - * Physical design (DM: physical)
 - * (Transaction design) (DM: process)

Conceptual design

- * Main idea: saving the data and retrieving it -> **DATA**
- * Information needed: everything you can find
- * Requirements and analysis usually half way -> need a lot of questions and answers
- * Interview the end users! Officially, unofficially
- * Model only the **target**, not the whole world


Conceptual design

- * Use right **terminology** and clear names, much easier to communicate with the end users (one of the reasons to model!)
- * Try to find and understand the **main** concepts and their relationships (these are the most difficult to change during the iterations)

Conceptual design, an example


- * "At one time one will make 1 or more Document/-s A and exactly one Document B for a Customer. One must know exactly which Document B was with which Document/-s A."
- * EASY AND CLEAR!

Conceptual design, an example




Conceptual design, an example

* OR?




Conceptual design, an example

* OR?




Conceptual design, an example

* OR?



Conceptual design, an example

* OR?



Conceptual design

- * Or one of the other dozen different ways to model this requirement?

Conceptual design

- * Modeling is **difficult** because
 - * Spoken/written language is not exact
 - * Usually all the "important" things are those that "everybody knows" so they are not told.
 - * At this stage we do not know one important thing: how the data will be **retrieved**? That will be on iteration 9...
- * Modeling is **mandatory** because when modeling the database designer realizes **what must be asked!**

Conceptual design, principles

- * Everybody has their own and preferred ways of modeling (“handwriting”) and that is ok, but it is important to be **consistent** at least inside one model or system or application.
- * Naming standards
- * Documentation
- * Comments from colleagues
- * Reviews, audits,...
- * ...

Design a Data Model for the Application

- * Not the "whole world"
- * System definition: the scope and boundaries of the database application (helps to define entities, attributes and relationships, understanding how this system sits to the whole architecture and other IT systems)
- * User views (helps to define roles and privileges, understanding how the system will be used)
- * database and application design are parallel activities of the database system development lifecycle

Designing the database

1. **Requirement analysis:** finding and analysing the requirements the future end users have

Result: specification of user requirements

- **data** requirements
- **functional** requirements

Also requirements for security, performance, ...

Designing the database

2. **Conceptual design.** "Interpretation" of all the requirements to a formal presentation (conceptual model).

Result: conceptual schema, also textual documentation is possible/recommended (to make sure all the knowledge is documented)

This is a tool for communication with end users.

Requirement analysis and Conceptual design

- * Collecting requirements and analyzing them
- * Fact-finding: interviews, questionnaires, existing documentation,... (recordings)
- * Requirements specifications
 - * data requirements
 - * functional requirements (performance, security, backup/recovery,..)
- * Completely neutral to any technology

Requirement analysis and Conceptual design

- * Why entity-relationship model (ER)
 - * Defining the tables directly based on requirements can be too difficult and lead to a wrong db schema.
 - * Based on a good ER it is easy to generate the relational model (which is at least on 3NF)
- * Data Flow Diagrams (DFD)

Methods

- * Collecting requirements
 - * Centralized: all requirements merged into a single set of requirements
 - * View: requirements for each user view as separate lists. Separate data models (local data model), that in later stage are merged as global data model.
 - * A combination of these two
- * Strategy
 - * bottom-up
 - * top-down
 - * Inside-out
 - * mixed

ER model, An Entity Type

- * A group of objects with the same properties
- * In spoken language: an Entity, a real life object ("A Customer", "An Invoice")
- * "A noun"

Strong and weak entity

- * strong entity: is not existence-dependent on some other entity (“Customer”, “Invoice”)
- * weak entity: is existence-dependent on another entity (“Invoiceline”)
 - * owner entity is identifying
 - * identifying relationship

ER model, An Attribute

- * Attribute is a property of an entity or a relationship ("The Firstname", "The OrderDate")

ER model, An Attribute

- * Atomic ("Firstname")
- * Composite attribute ("Name" = "Firstname"+"Lastname")
- * Multivalued/set-valued ("Phonenumber")
- * Derived attribute ("Age", "Duration")
- * A Key ("StudentID")
- * Mandatory/non-mandatory (NULLs)

ER model, A Relationship Type

- * In spoken language: a relationship
- * A Relationship between 2 or more entities
- * Entities *participate* the relationship
- * An Entity can participate a relationship in different Roles (Recursive Relationship)
- * Degree of a relationship is the number of participating entities in a relationship.
- * A relationship can have its own Attributes
- * "A verb"

ER model, A Relationship

- * Constraints
 - * **Cardinality** constraint
 - * 1:1
 - * 1:m
 - * m:1
 - * m:n
 - * **Participation** constraint
 - * Total/partial (mandatory/optional)

ER model, A Key

- * A **Candidate Key**: The minimal set of attributes that uniquely identifies each occurrence of an entity.
- * Composite key is a candidate key that consists of two or more attributes.
- * The **Primary Key**: one selected Candidate Key to identify each occurrence of an entity. Other possible candidate keys are called **alternate keys**.
- * Which one to choose? (will it stay unique, not changing, data type, length,...)

ER model, A Key

- * The Primary Key
 - * Natural Key
 - * Surrogate Key
 - * A combination of these two...
- * Not a required feature for an ER model, but relational model needs a Primary Key.

ER model, A Key

- * Natural Key usually more efficient because joins are sometimes not needed at all
- * A surrogate key is not often used in retrieving data (because it has no meaning to the business), but PK always has an index -> wasted space
- * Surrogate does not stop inserting logical duplicates (UK needed)

Designing

Avoid Redundance

- * A wrong place for an attribute will cause redundancy
 - * For example Country in the address entity
- * Too many relationships may cause redundancy

Normalization

- * Normalization is a method that **removes data redundancy** from a relation (minimizing the insertion, deletion and update anomalies that degrade the performance of databases)
- * Normalization uses a series of tests (normal forms) to help identify the **optimal grouping for these attributes** to identify a set of suitable relations that supports the data requirements.
- * examining the **functional dependencies** between attribute

Functional Dependency

- * a relationship between two attributes in a relation (determinant, dependent)
 - * ISBN \rightarrow BookTitle (ISBN determines the book title)
 - * or
 - * Employee_id \rightarrow name job salary
 - * or
 - * street city \rightarrow postcode and
 - * postcode \rightarrow city

Armstrong's axioms

- * **Axiom of reflexivity**

- * If $Y \subseteq X$, then $X \rightarrow Y$
- * (If Y is a subset of X , then X determines Y .)

- * **Axiom of augmentation**

- * If $X \rightarrow Y$, then $XZ \rightarrow YZ$
- * (If X determines Y , then XZ determines YZ .)

- * **Axiom of transitivity**

- * If $X \rightarrow Y$ ja $Y \rightarrow Z$, then $X \rightarrow Z$
- * (If X determines Y and Y determines Z , then X must also determine Z .)

Why Functional Dependency?

- * To define keys
 - * X = attributes never on the right hand side
 - * Y = attributes both on the right hand and the left hand side
 - * Z = attributes on the right hand never on the left hand side
 - * X must be in every key, Y might be in a key, Z cannot be in any key
- * To normalize

- * `address(street, postcode, city)`
- * `address1(street, postcode)` and
- * `address2(postcode, city)`
- * inclusion dependency
 - * $address1[postcode] \subseteq address2[postcode]$

Normalization and the “Forms”

- * First Normal Form (1NF)
 - * none of its domains have any sets as elements, attributes contains only single values from the domain (each column contains atomic values, and there are not repeating groups of columns)
- * Second Normal Form (2NF)
 - * 1NF+all non-key attributes are fully functional dependent on the primary key (the columns are dependent on the primary key)

Normalization and the “Forms”

- * Third Normal Form (3NF)
 - * 2NF+there is **no** transitive functional dependency (the columns are dependent on the primary key (2NF) and no other columns in the table)
 - * Student_id, Student_name, Street, City, State, Zip
 - * (Student_id as primary key)
 - * But now Street, City, State are dependent on Zip not Student_id -> transitive functional dependency -> not in 3NF

Normalization and the “Forms”


- * Boyce–Codd Normal Form (BCNF)
 - * Must be in 3NF and for each functional dependency ($X \rightarrow Y$), X should be a super Key.
 - * if and only if there are no non-trivial functional dependencies of attributes on anything other than a superset of a candidate key. (Each attribute must represent a fact about the key, the whole key, and nothing but the key)
- * Fourth Normal Form (4NF), Fifth Normal Form (5NF).

An entity, a relationship or an attribute?

- * There are no foreign keys (and FK columns) in ER!
 - * They are **relationships**
- * An entity or an attribute?
 - * Example, phonenumber (how many? Other attributes?)
- * An entity or a relationship?
 - * If it is a relationship in real world it is most likely a relationship in the data model too.

Specialization

- * Specialization: A customer -> Person or Company




Specialization

- * Can Person change to a Company or vice versa?
- * Do you always know which one the customer is?
- * ...


Generalization

- * Generalization: A Vehicle (<- a car, a bus, a lorry, ...)
 - * A bus might have some attributes a car does not have etc. NULLs.
 - * The Customer example:



Generalization

- * Usually a discriminator to identify the "type"



- * Now only the attributes all the "types" have can be defined mandatory and there will be plenty on NULLs

Specialization, generalization

- * **Participation** constraint

- * Defines whether every member in the superclass must also be a member of a subclass or not
- * mandatory or optional

- * **Disjoint** constraints

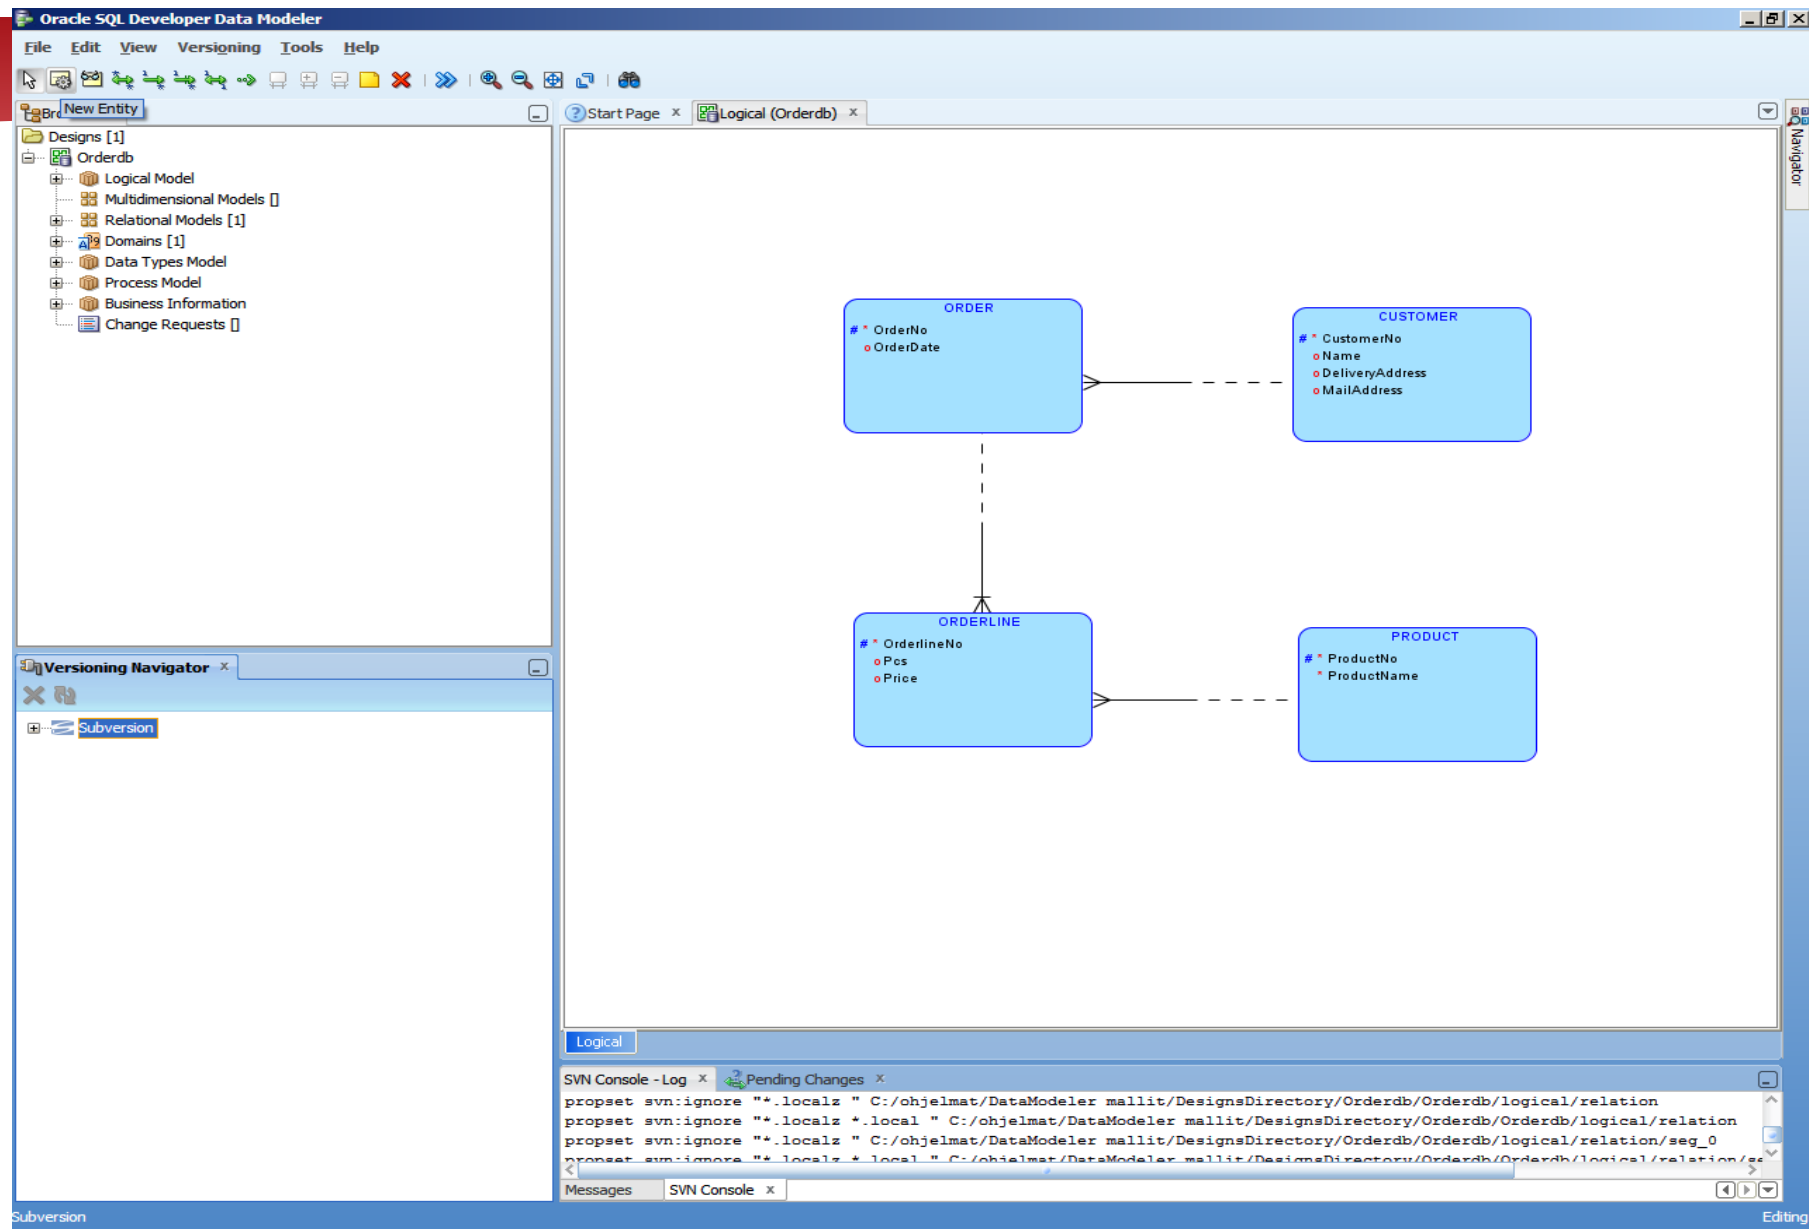
- * The disjoint constraint only applies when a superclass has more than one subclass. If the subclasses are disjoint, then an entity occurrence can be a member of only one of the subclasses.

Data Modeler

- * Let's design a database
- * <https://github.com/Mavihe1/DBworkshop>



Testing the conceptual model



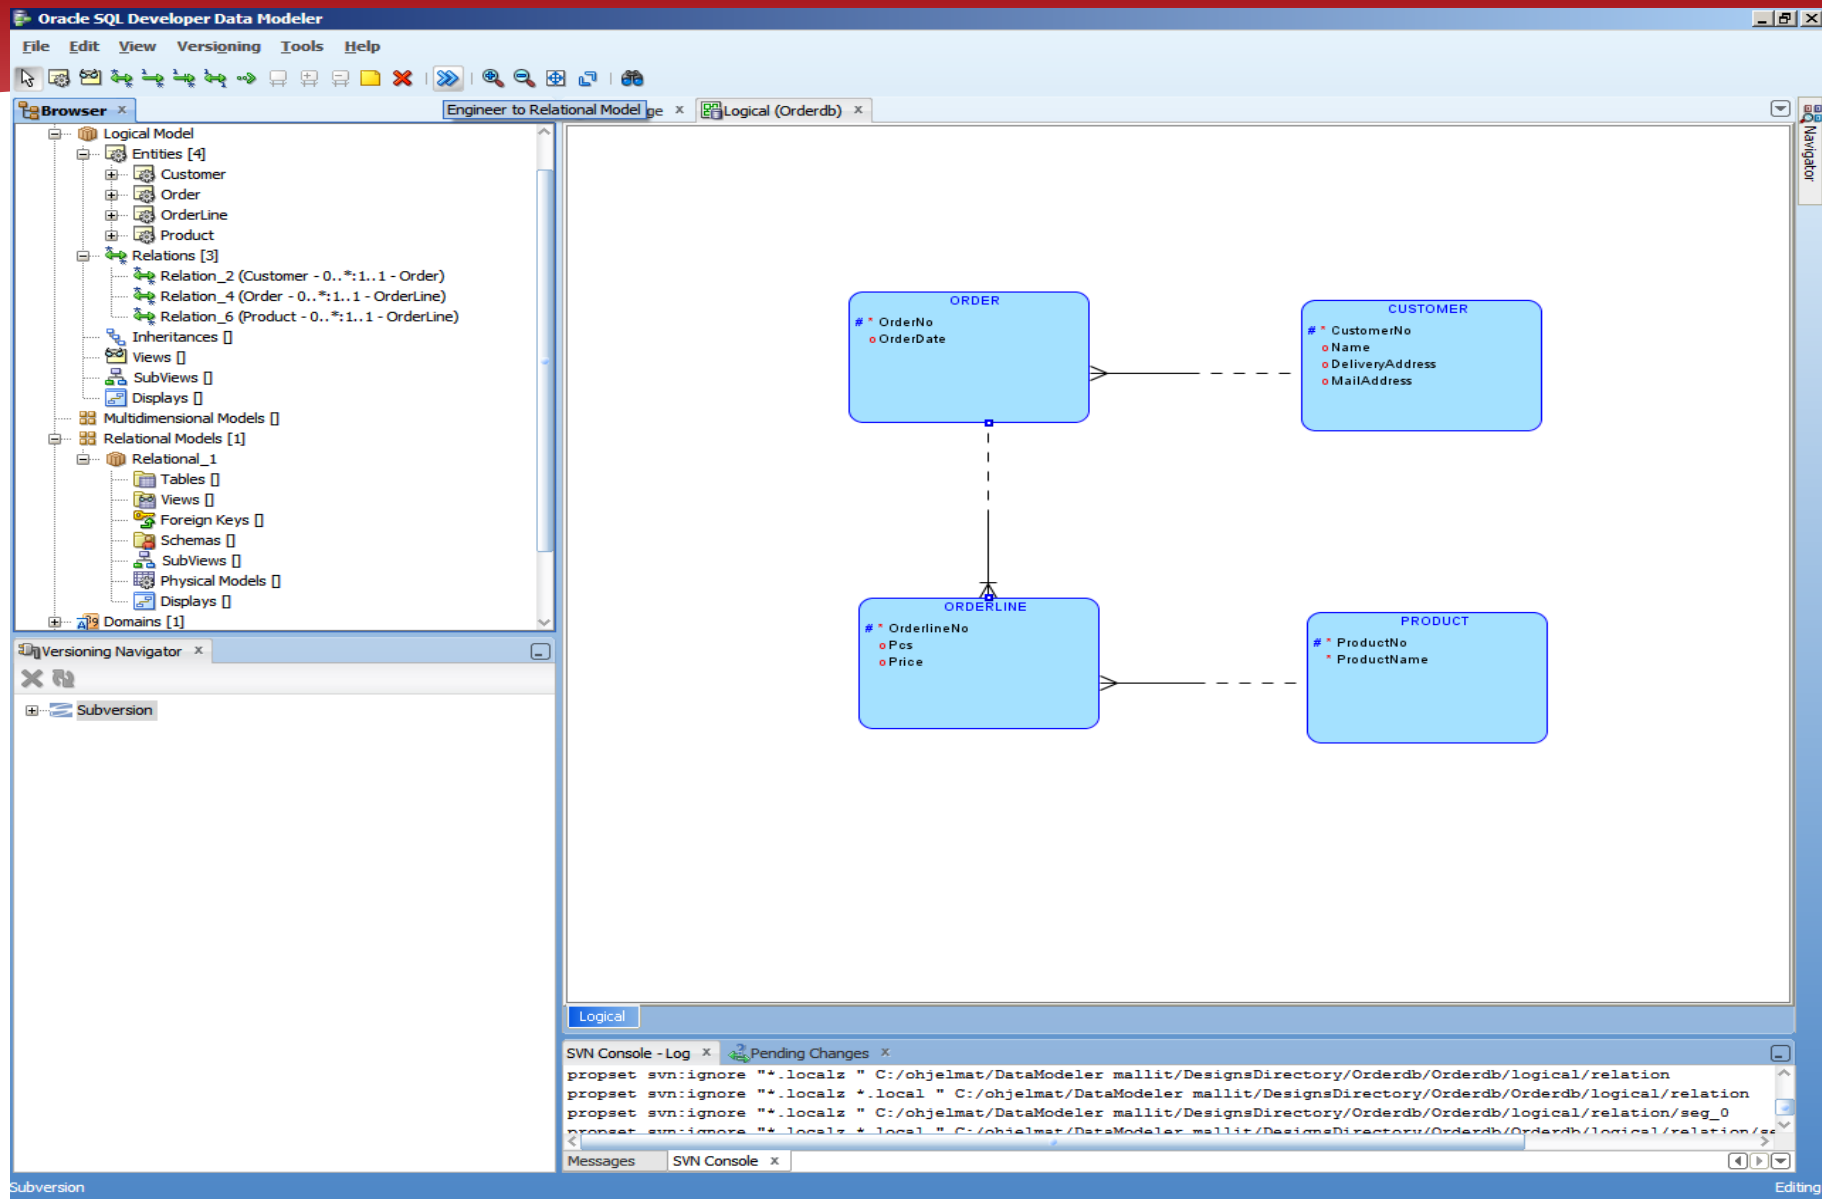
Designing the database

3. **Logical design:** transforming the conceptual model into a logical data model and a logical schema that the RDMS understands

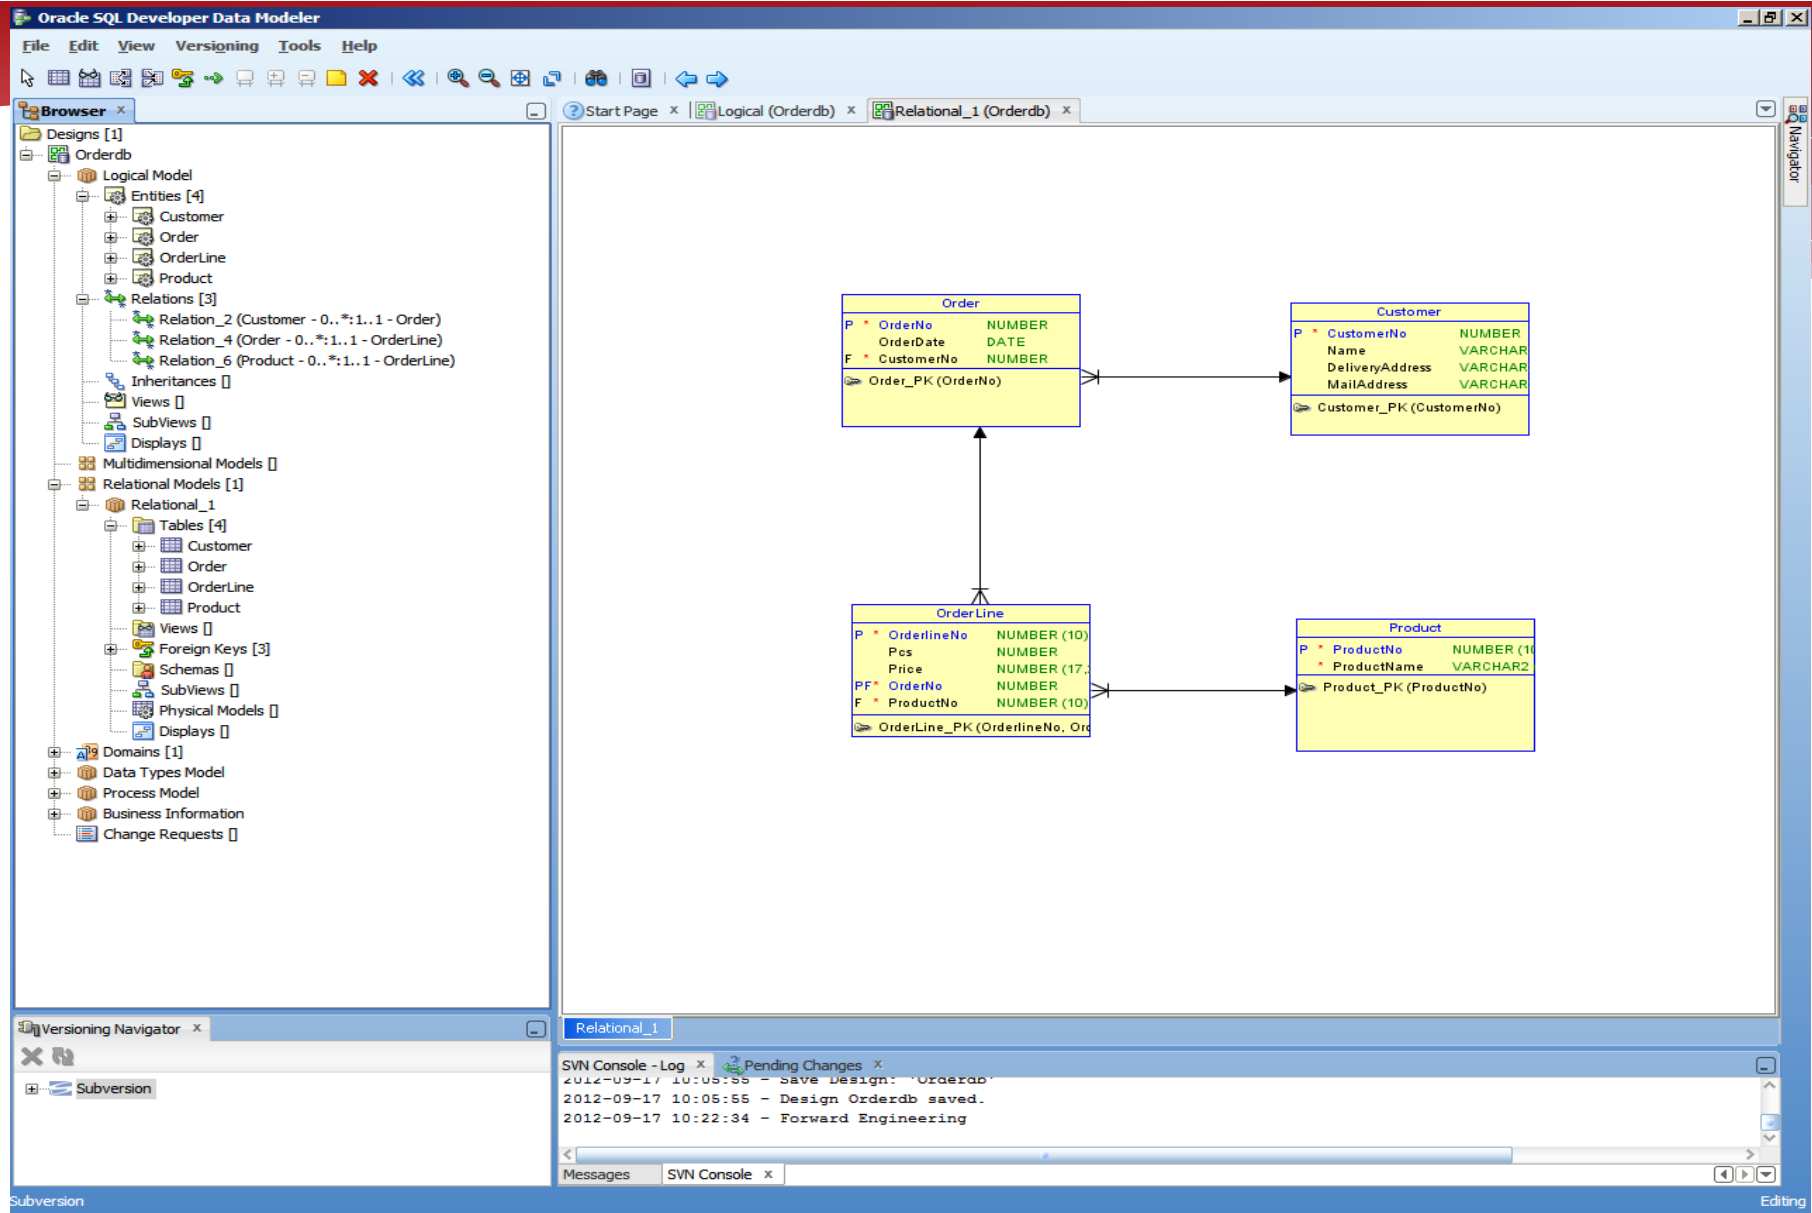
Result: relational-database schema
(relational schemas and constraints)

To logical (relational) model

- * The conceptual data model is refined and mapped to a logical data model.
- * The logical data model defines the DBMS used
 - * In our case relational data model for RDBMS
- * But we do not define the "brand" yet (Oracle, MySQL, MS SQL Server,...)
- * The easiest is if you have a tool that automatically creates tables and foreign keys based on the ER



Relational model



Designing the database

4. **Physical design:** instances, tablespaces, indexes, disks ...

And all of these phases over and over again...

(this is different, we have always done that but not so many times and in such a short cycles)

From relational to physical

- * how the database will be implemented
- * A specific DBMS system

Physical design

- * In short:
 - * Start with the logical model
 - * Design the disc usage, count the need for the space etc.
 - * Indexes
 - * Schemas
 - * Tablespaces, storage definitions, user permissions,...

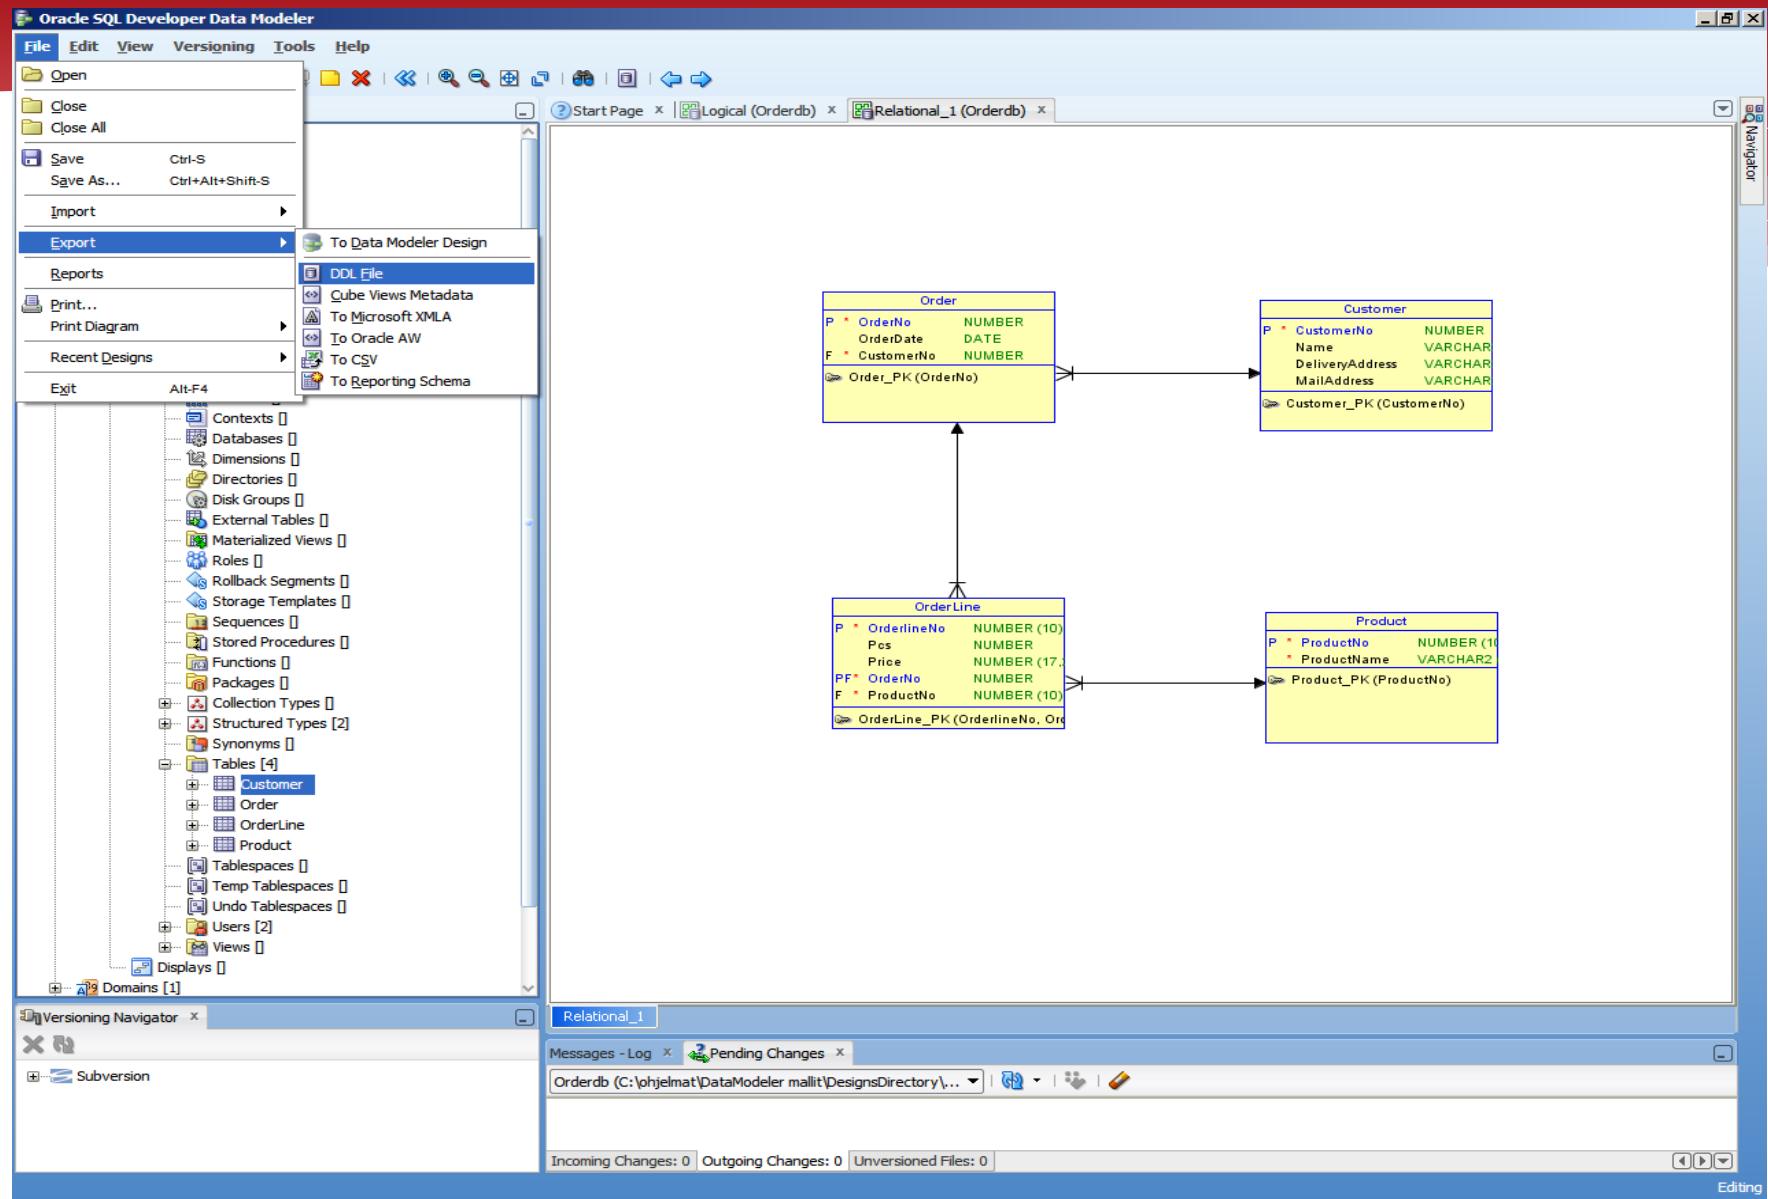
Physical design


- * How, who and where to document?
- * How to document any changes (patch, changes in database objects, ...)?
- * Backup and recovery (test that it works every now and then!)


Physical model



Implementation: Physical to DDLs







DDL File Editor - Oracle Database 11g

Oracle Database 11g Generate Clear Find

```
-- Generated by Oracle SQL Developer Data Modeler 3.1.0.683
-- at:      2011-12-13 13:11:12 EET
-- site:    Oracle Database 11g
-- type:    Oracle Database 11g

CREATE TABLE Customer
(
  CustomerNo NUMBER NOT NULL ,
  name VARCHAR2 (50) NOT NULL ,
  DeliveryAddress VARCHAR2 (200) ,
  MailAddress VARCHAR2 (50)
) LOGGING
;

ALTER TABLE Customer
  ADD CONSTRAINT "Customer PK" PRIMARY KEY ( CustomerNo ) ;

CREATE TABLE "Order"
(
  OrderNo NUMBER NOT NULL ,
  OrderDate DATE NOT NULL ,
  CustomerNo NUMBER
) LOGGING
;

ALTER TABLE "Order"
  ADD CONSTRAINT "Order PK" PRIMARY KEY ( OrderNo ) ;
```

Save Close Help

That was the process... but

What kind of DB are you designing?

- * Different rules for OLTP and DW
 - * My "own" data
 - * Somebody else's data

Maintenance

- * Monitoring the performance of the system: tuning if needed.
- * Maintaining and upgrading the database system.
- * Maintaining the user privileges.
- * Testing backups and recoveries.
- * New requirements implemented using the design process.

Performance

- * Plan, count, ...
- * Test! With a real material (size, quality,...) and the right versions of program.

Tuning

- * Hardware and operating system
- * DBMS
- * Model level (logical, physical, transaction)
- * All three levels affect each others

Indexing

- * Select **might** be faster
- * Insert, Update, Delete slower
- * Not a silver bullet

Tuning the queries

- * Transform the query to a more efficient form (**execution plan**) with the same outcome (**result**) than the original query (**equivalent**).

Tuning, the process

- * Ask direct questions. ("Nothing is working" is not an answer)
- * Be systematic and find the real problem
- * Solve the problem
- * Document the problem and the solution (new problems because of the solution?)

Refactoring the database

- * Analyze the change from the **database** perspective
 - * The phase of the database, development vs production
 - * The phase of usage of the database, development vs production
 - * Completely new database vs "old" database
- * Analyze the change from the **project** perspective (cheep coding now is not an excuse for a bad database design)

Refactoring the database

- * Analyze the change from the **maintenance** point of view (would it be more clear to call discount DISCOUNT than AMOUNT?)
- * Versioning
- * Always remember somebody will maintain this system and this database
- * Always remember clear and understandable solutions, do not try to be too clever

Refactoring the database

- * The problems related to refactoring the database usually have nothing to do with the database but everything to do "on top" of the database
- * Refactoring the database is nothing special, just normal database design process

A tool

Oracle SQL Developer Data Modeler

- * To be efficient in designing you need a tool: my recommendation is Data Modeler
 - * Free of charge
 - * Support for many different databases
 - * Support for both documenting the existing databases and designing a new one (and maintaining that)
 - * Support for reporting, naming standards, glossaries, design rules, ...
 - * Support for version control and multiuser environment
 - * Support for everything you need for database design plus more

Conclusions

- * Database designing is an important work
- * The database must be designed and for the right purpose .
- * Data integrity and data quality must be the guiding rules.
- * 3NF or Boyce-Codd preferred

Conclusions

- * You need a tool!
- * Data Modeler is a good tool; good support for iterative processes
- * Enables documenting and versioning (and comparing the versions)
- * Enables multiuser environment
- * Is free to use
- * Support for other databases as well

THANK YOU!

QUESTIONS?

heli@miracleoy.fi

Twitter: @HeliFromFinland

Blog: Helifromfinland.com